


Protection: Address Spaces


<http://inst.eecs.berkeley.edu/~cs162>

Goals for Today

- Kernel vs User Mode 
- What is an Address Space?
- How is it Implemented?
- Discussion of Dual-Mode operation
- Comparison among options

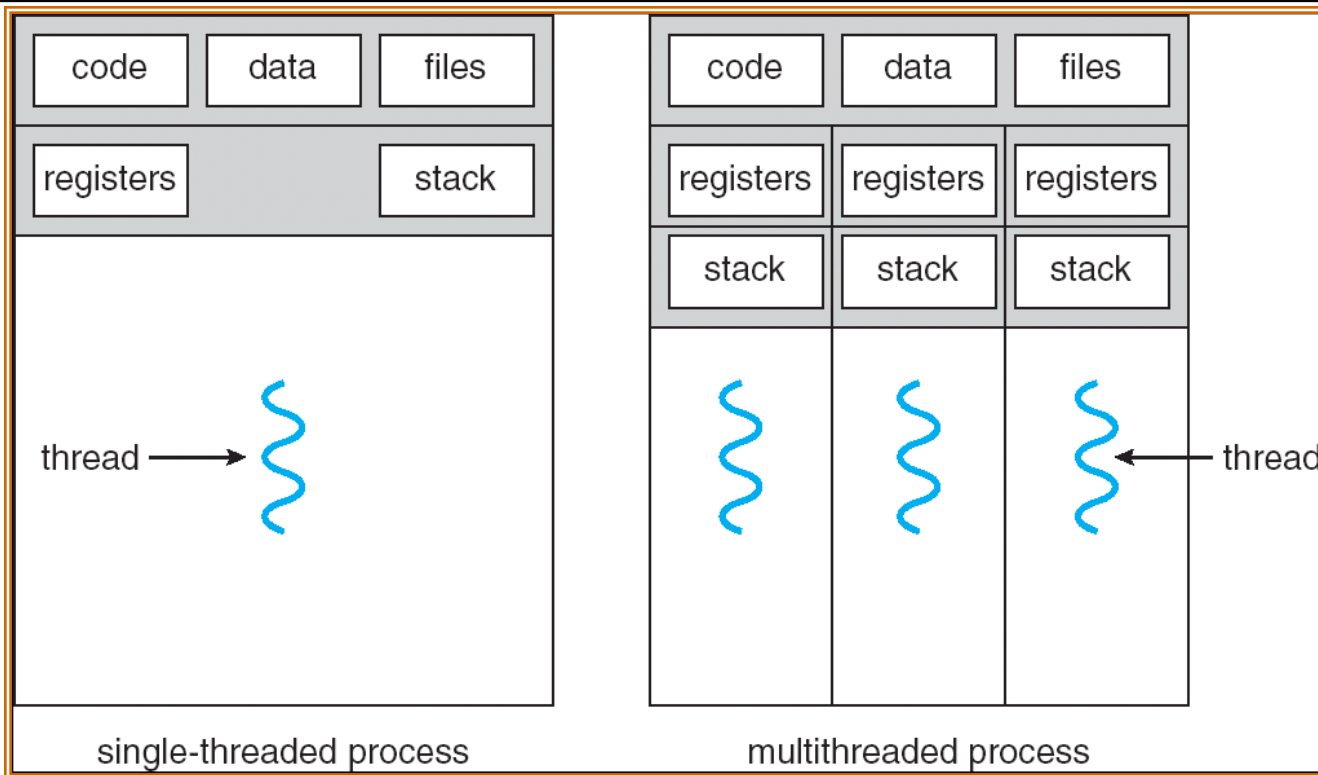
Note: Some slides and/or pictures in the following are adapted from slides ©2005 Silberschatz, Galvin, and Gagne

Virtualizing Resources



- **Physical Reality:**
Different Processes/Threads share the same hardware
 - Need to multiplex CPU (Just finished: scheduling)
 - Need to multiplex use of Memory (Today)
 - Need to multiplex disk and devices (later in term)
- **Why worry about memory sharing?**
 - The complete working state of a process and/or kernel is defined by its data in memory (and registers)
 - Consequently, cannot just let different threads of control use the same memory
 - » Physics: two different pieces of data cannot occupy the same locations in memory
 - Probably don't want different threads to even have access to each other's memory (protection)

Recall: Single and Multithreaded Processes



- **Threads encapsulate concurrency**
 - "Active" component of a process
- **Address spaces encapsulate protection**
 - Keeps buggy program from trashing the system
 - "Passive" component of a process

Important Aspects of Memory Multiplexing

- **Controlled overlap:**

- Separate state of threads should not collide in physical memory. Obviously, unexpected overlap causes chaos!
- Conversely, would like the ability to overlap when desired (for communication)

- **Translation:**

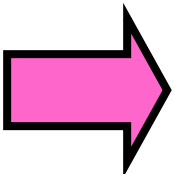
- Ability to translate accesses from one address space (virtual) to a different one (physical)
- When translation exists, processor uses virtual addresses, physical memory uses physical addresses
- Side effects:
 - » Can be used to avoid overlap
 - » Can be used to give uniform view of memory to programs

- **Protection:**

- Prevent access to private memory of other processes
 - » Different pages of memory can be given special behavior (Read Only, Invisible to user programs, etc).
 - » Kernel data protected from User programs
 - » Programs protected from themselves

Binding of Instructions and Data to Memory

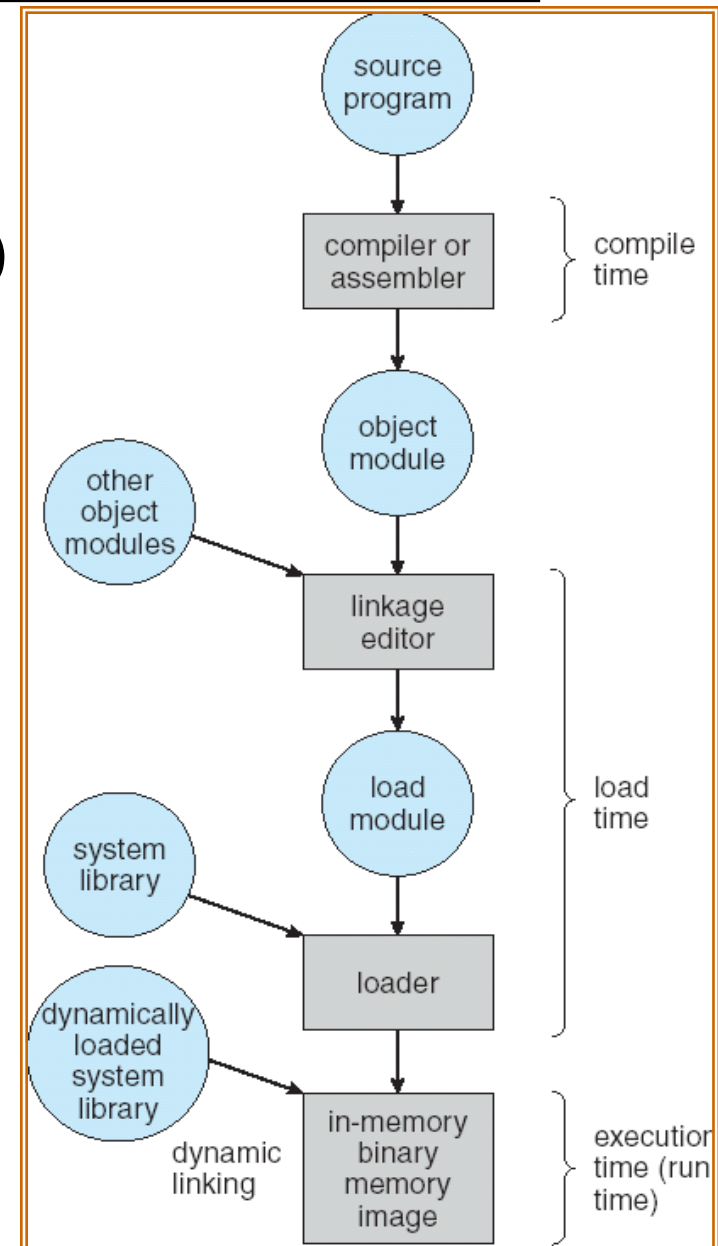
- Binding of instructions and data to addresses:
 - Choose addresses for instructions and data from the standpoint of the processor

data1: dw 32		0x300 00000020
...		...
start: lw r1, 0(data1)		0x900 8C2000C0
jal checkit		0x904 0C000340
loop: addi r1, r1, -1		0x908 2021FFFF
bnz r1, r0, loop		0x90C 1420FFFF
...		...
checkit: ...		0xD00 ...

- Could we place data1, start, and/or checkit at different addresses?
 - » Yes
 - » When? Compile time/Load time/Execution time
- Related: which physical memory locations hold particular instructions or data?

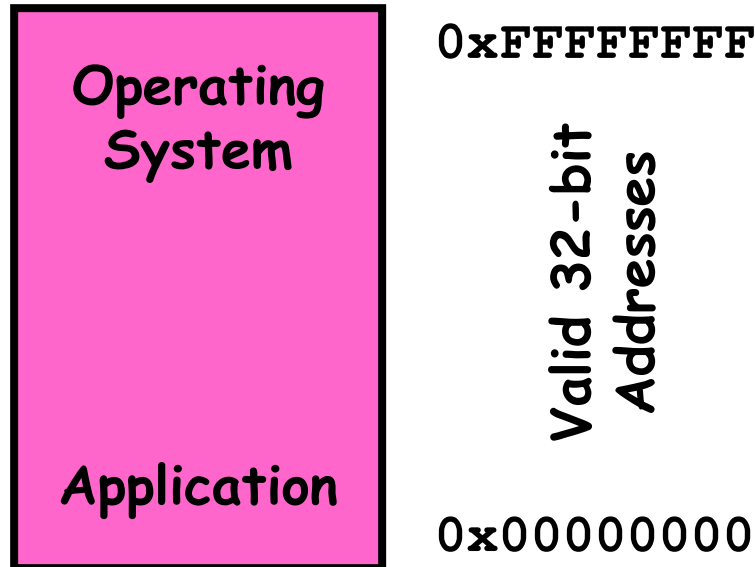
Multi-step Processing of a Program for Execution

- Preparation of a program for execution involves components at:
 - Compile time (i.e. "gcc")
 - Link/Load time (unix "ld" does link)
 - Execution time (e.g. dynamic libs)
- Addresses can be bound to final values anywhere in this path
 - Depends on hardware support
 - Also depends on operating system
- Dynamic Libraries
 - Linking postponed until execution
 - Small piece of code, *stub*, used to locate the appropriate memory-resident library routine
 - Stub replaces itself with the address of the routine, and executes routine



Recall: Uniprogramming

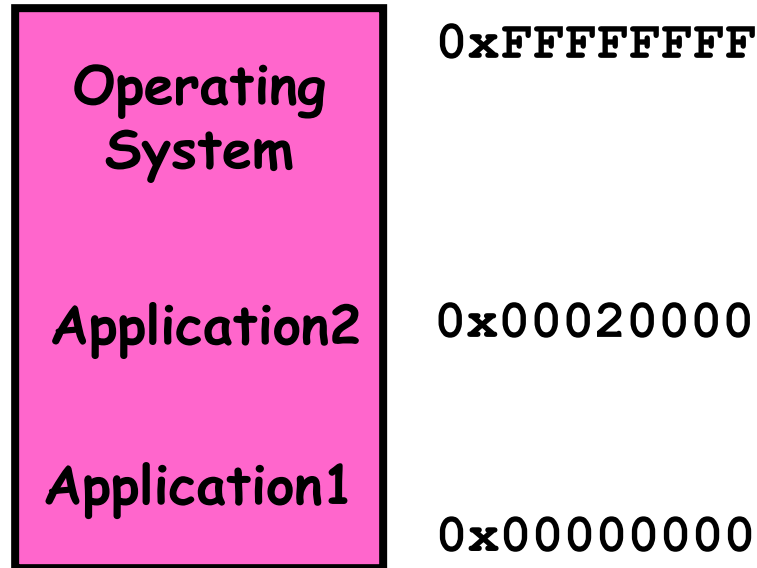
- Uniprogramming (no Translation or Protection)
 - Application always runs at same place in physical memory since only one application at a time
 - Application can access any physical address



- Application given illusion of dedicated machine by giving it reality of a dedicated machine
- Of course, this doesn't help us with multithreading

Multiprogramming (First Version)

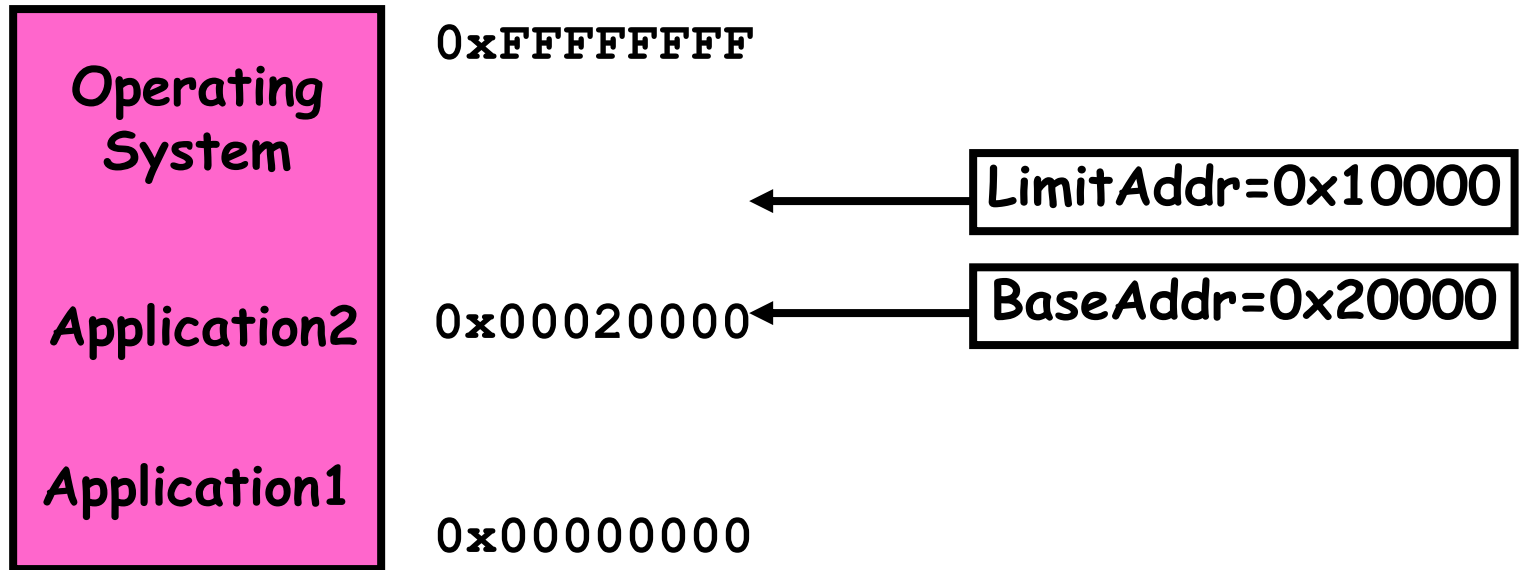
- Multiprogramming without Translation or Protection
 - Must somehow prevent address overlap between threads



- Trick: Use Loader/Linker: Adjust addresses while program loaded into memory (loads, stores, jumps)
 - » Everything adjusted to memory location of program
 - » Translation done by a linker-loader
 - » Was pretty common in early days
- With this solution, no protection: bugs in any program can cause other programs to crash or even the OS


Multiprogramming (Version with Protection)

- Can we protect programs from each other without translation?



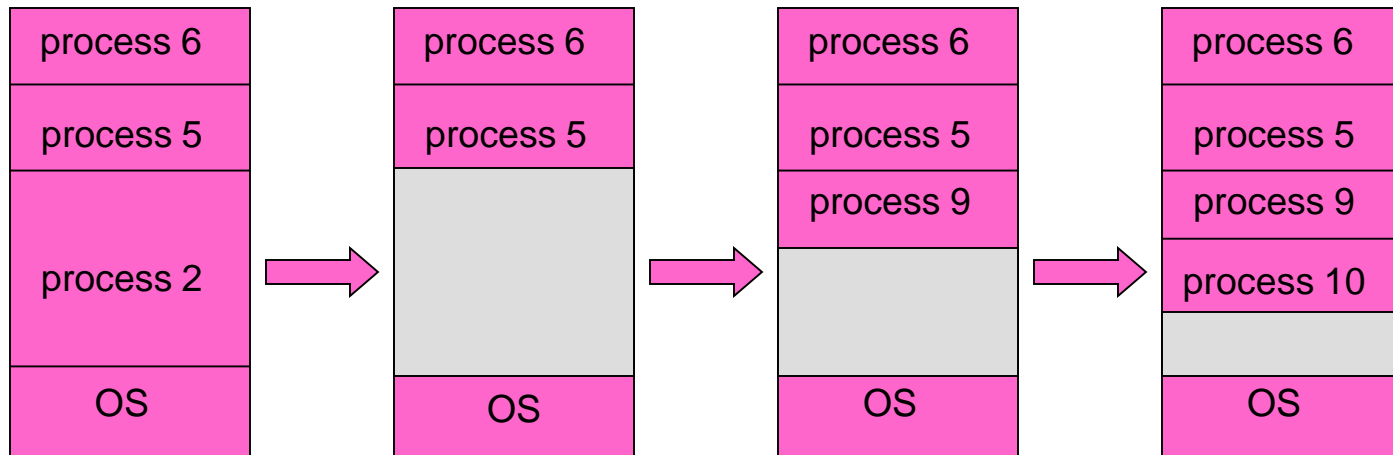
- Yes: use two special registers *BaseAddr* and *LimitAddr* to prevent user from straying outside designated area
 - » If user tries to access an illegal address, cause an error
- During switch, kernel loads new base/limit from TCB
 - » User not allowed to change base/limit registers

Segmentation with Base and Limit registers



- Could use base/limit for **dynamic address translation** (often called "segmentation"):
 - Alter address of every load/store by adding "base"
 - User allowed to read/write within segment
 - » Accesses are relative to segment so don't have to be relocated when program moved to different segment
 - User may have multiple segments available (e.g x86)
 - » Loads and stores include segment ID in opcode:
x86 Example: `mov [es:bx], ax.`
 - » Operating system moves around segment base pointers as necessary

Issues with simple segmentation method

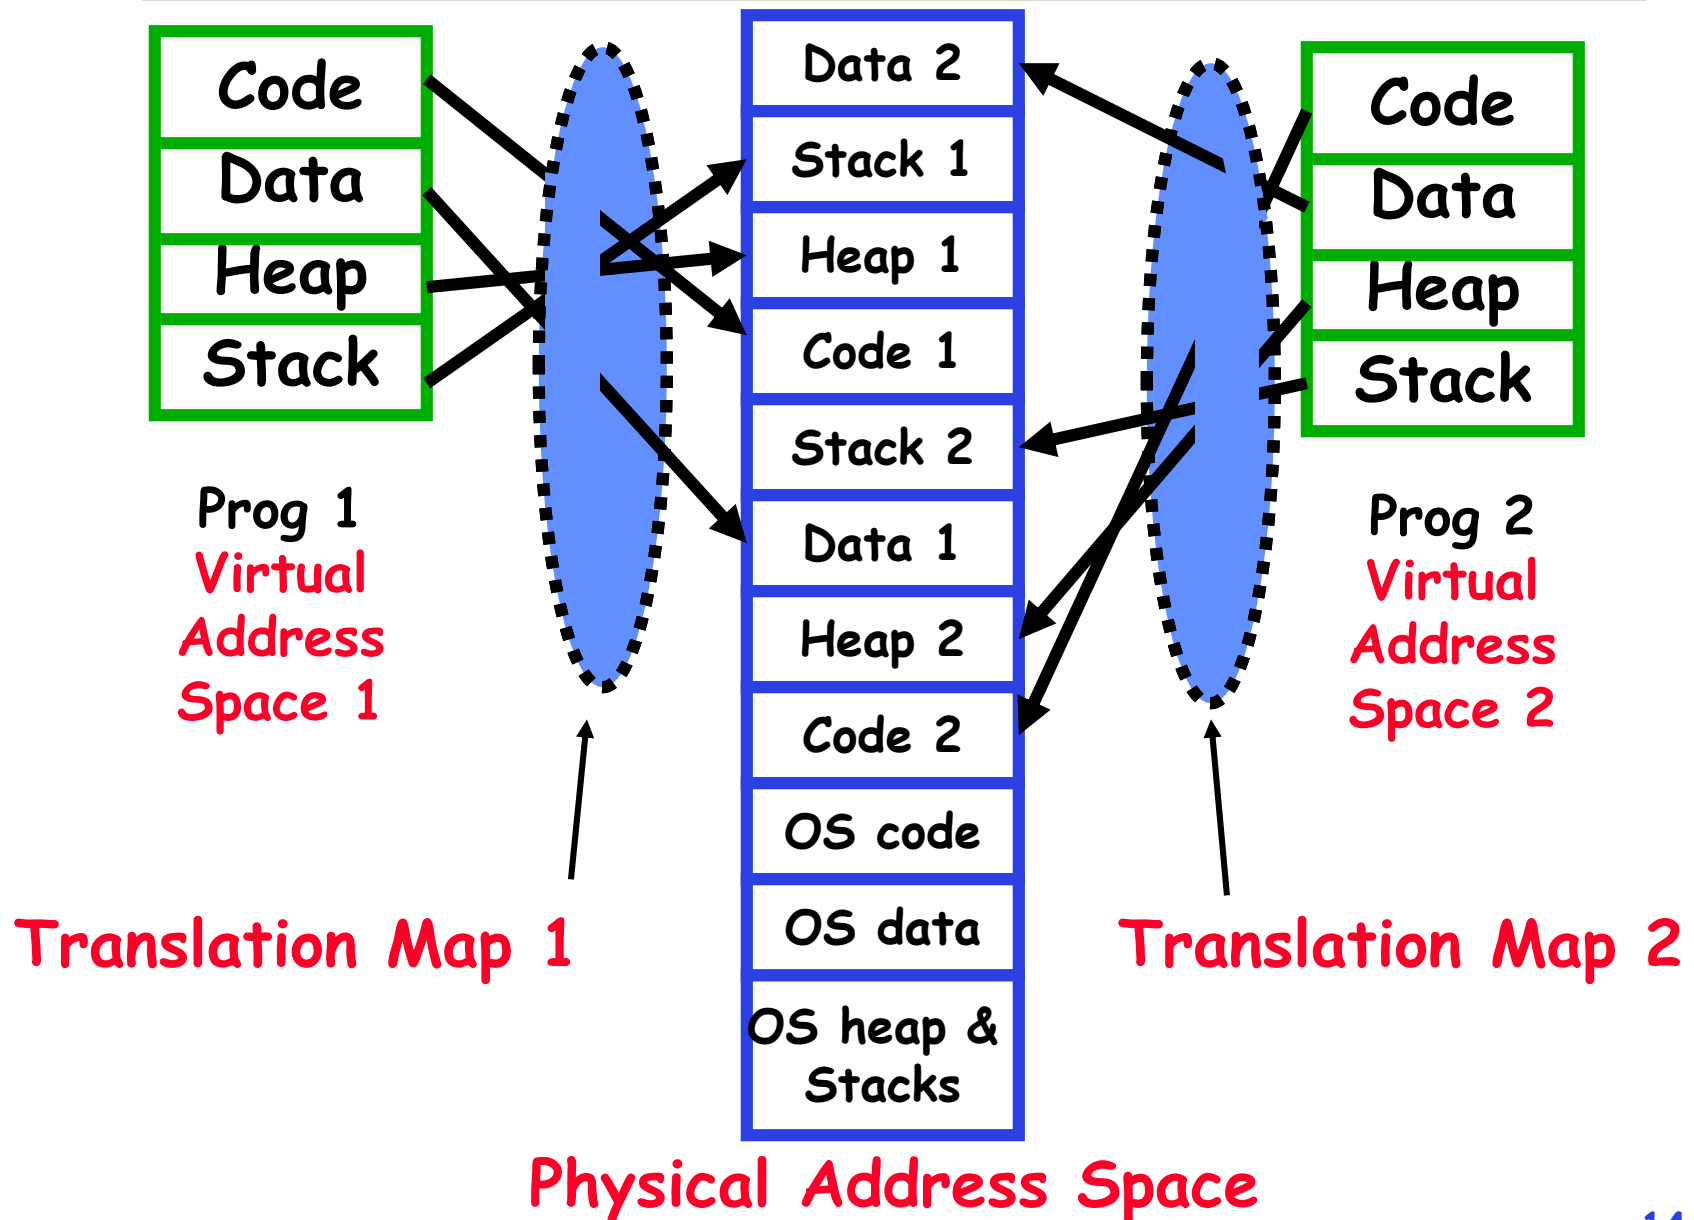


- **Fragmentation problem**
 - Not every process is the same size
 - Over time, memory space becomes fragmented
- **Hard to do inter-process sharing**
 - Want to share code segments when possible
 - Want to share memory between processes
 - Helped by providing multiple segments per process
- **Need enough physical memory for every process**


Multiprogramming (Translation and Protection version 2)

- Problem: Run multiple applications in such a way that they are protected from one another
- Goals:
 - Isolate processes and kernel from one another
 - Allow flexible translation that:
 - » Doesn't lead to fragmentation
 - » Allows easy sharing between processes
 - » Allows only part of process to be resident in physical memory
- (Some of the required) Hardware Mechanisms:
 - General Address Translation
 - » Flexible: Can fit physical chunks of memory into arbitrary places in users address space
 - » Not limited to small number of segments
 - » Think of this as providing a large number (thousands) of fixed-sized segments (called "pages")
 - Dual Mode Operation
 - » Protection base involving kernel/user distinction

Example of General Address Translation



Two Views of Memory

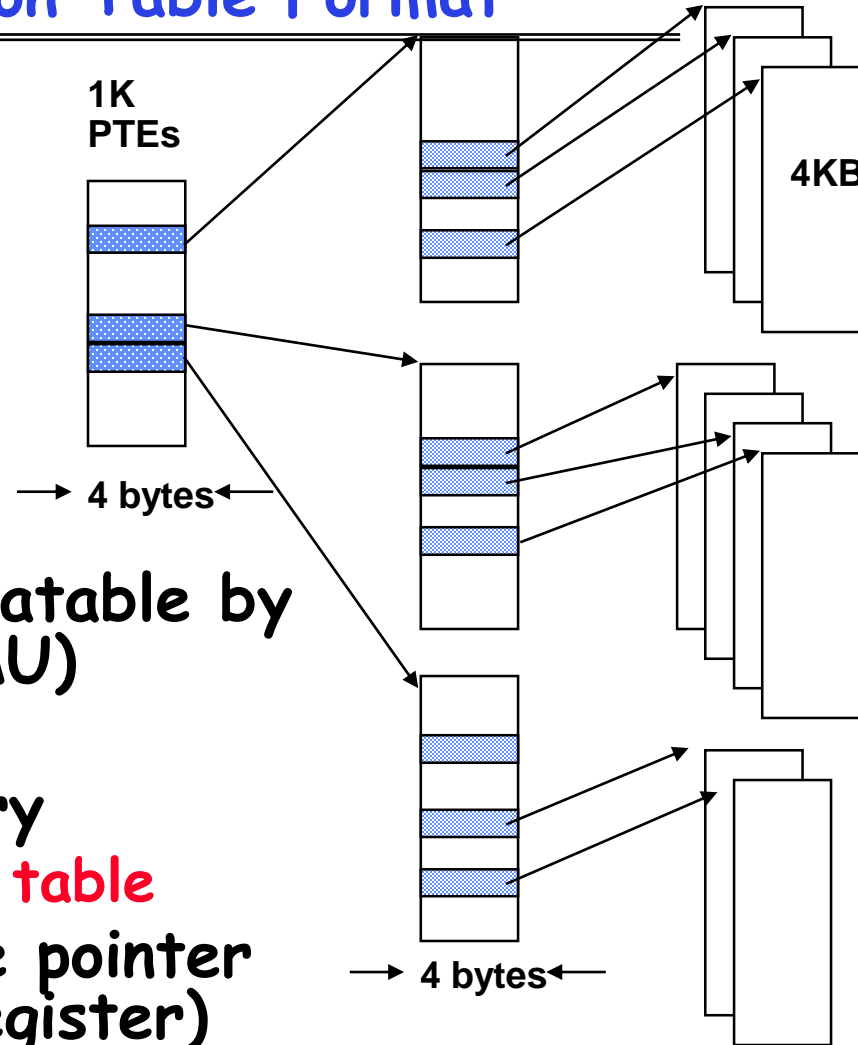


- **Recall: Address Space:**
 - All the addresses and state a process can touch
 - Each process and kernel has different address space
- **Consequently: two views of memory:**
 - View from the CPU (what program sees, virtual memory)
 - View from memory (physical memory)
 - Translation box converts between the two views
- **Translation helps to implement protection**
 - If task A cannot even gain access to task B's data, no way for A to adversely affect B
- **With translation, every program can be linked/loaded into same region of user address space**
 - Overlap avoided through translation, not relocation

Example of Translation Table Format

Two-level Page Tables

32-bit address:



- Page: a unit of memory translatable by memory management unit (MMU)
 - Typically 1K - 8K
- Page table structure in memory
 - Each user has different page table
- Address Space switch: change pointer to base of table (hardware register)
 - Hardware traverses page table (for many architectures)
 - MIPS uses software to traverse table

Dual-Mode Operation

- Can Application Modify its own translation tables?
 - If it could, could get access to all of physical memory
 - Has to be restricted somehow
- To Assist with Protection, **Hardware** provides at least two modes (Dual-Mode Operation):
 - "Kernel" mode (or "supervisor" or "protected")
 - "User" mode (Normal program mode)
 - Mode set with bits in special control register only accessible in kernel-mode
- Intel processor actually has four "rings" of protection:
 - PL (Privilege Level) from 0 - 3
 - » PLO has full access, PL3 has least
 - Privilege Level set in code segment descriptor (CS)
 - Mirrored "IOPL" bits in condition register gives permission to programs to use the I/O instructions
 - Typical OS kernels on Intel processors only use PLO ("user") and PL3 ("kernel")

For Protection, Lock User-Programs in Asylum

- **Idea: Lock user programs in padded cell with no exit or sharp objects**
 - Cannot change mode to kernel mode
 - User cannot modify page table mapping
 - Limited access to memory: cannot adversely effect other processes
 - » Side-effect: Limited access to memory-mapped I/O operations (I/O that occurs by reading/writing memory locations)
 - Limited access to interrupt controller
 - What else needs to be protected?
- **A couple of issues**
 - How to share CPU between kernel and user programs?
 - » Kinda like both the inmates and the warden in asylum are the same person. How do you manage this???
 - How do programs interact?
 - How does one switch between kernel and user modes?
 - » OS → user (kernel → user mode): getting into cell
 - » User → OS (user → kernel mode): getting out of cell

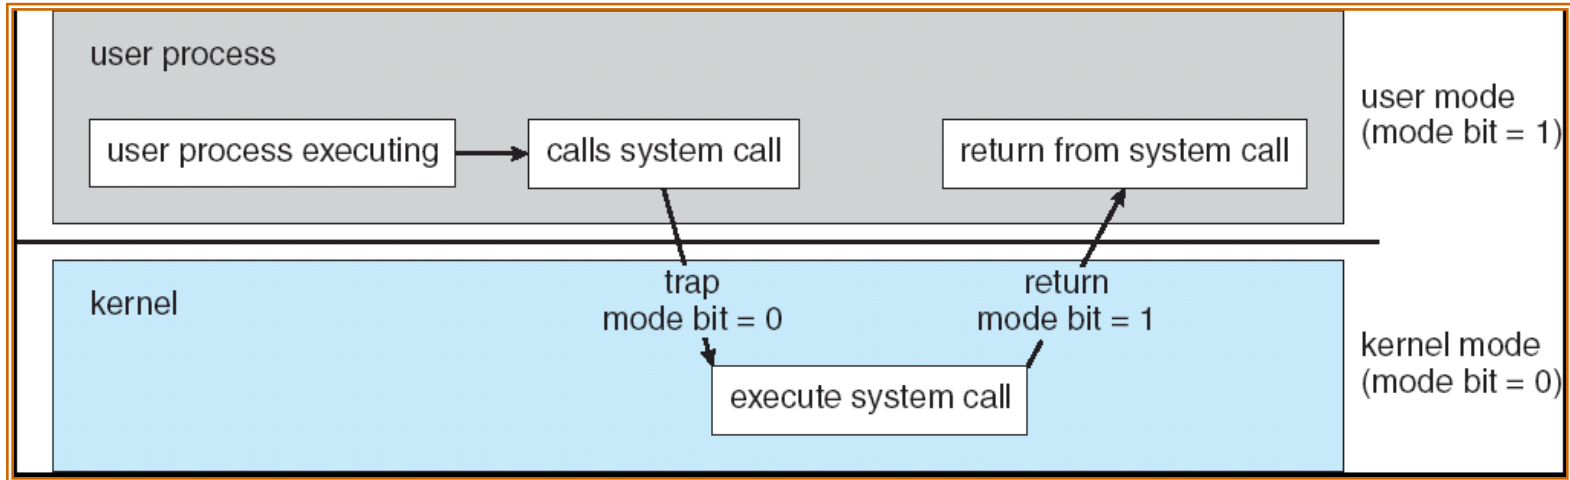


How to get from Kernel→User

- What does the kernel do to create a new user process?
 - Allocate and initialize address-space control block
 - Read program off disk and store in memory
 - Allocate and initialize translation table
 - » Point at code in memory so program can execute
 - » Possibly point at statically initialized data
 - Run Program:
 - » Set machine registers
 - » Set hardware pointer to translation table
 - » Set processor status word for user mode
 - » Jump to start of program
- How does kernel switch between processes?
 - Same saving/restoring of registers as before
 - Save/restore PSL (hardware pointer to translation table)

User→Kernel (System Call)

- Can't let inmate (user) get out of padded cell on own
 - Would defeat purpose of protection!
 - So, how does the user program get back into kernel?



- **System call:** Voluntary procedure call into kernel
 - Hardware for controlled User→Kernel transition
 - Can any kernel routine be called?
 - » No! Only specific ones.
 - System call ID encoded into system call instruction
 - » Index forces well-defined interface with kernel

System Call Continued


- What are some system calls?
 - I/O: open, close, read, write, lseek
 - Files: delete, mkdir, rmdir, truncate, chown, chgrp, ..
 - Process: fork, exit, wait (like join)
 - Network: socket create, set options
- Are system calls constant across operating systems?
 - Not entirely, but there are lots of commonalities
 - Also some standardization attempts (POSIX)
- What happens at beginning of system call?
 - » On entry to kernel, sets system to kernel mode
 - » Handler address fetched from table/Handler started
- System Call argument passing:
 - In registers (not very much can be passed)
 - Write into user memory, kernel copies into kernel mem
 - » User addresses must be translated!w
 - » Kernel has different view of memory than user
 - Every Argument must be explicitly checked!

User→Kernel (Exceptions: Traps and Interrupts)

- A system call instruction causes a synchronous exception (or “trap”)
 - In fact, often called a software “trap” instruction
- Other sources of *Synchronous Exceptions*:
 - Divide by zero, Illegal instruction, Bus error (bad address, e.g. unaligned access)
 - Segmentation Fault (address out of range)
 - Page Fault (for illusion of infinite-sized memory)
- Interrupts are *Asynchronous Exceptions*
 - Examples: timer, disk ready, network, etc....
 - **Interrupts can be disabled, traps cannot!**
- On system call, exception, or interrupt:
 - Hardware enters kernel mode with interrupts disabled
 - Saves PC, then jumps to appropriate handler in kernel
 - For some processors (x86), processor also saves registers, changes stack, etc.
- Actual handler typically saves registers, other CPU state, and switches to kernel stack

Additions to MIPS ISA to support Exceptions?

- Exception state is kept in “Coprocessor 0”
 - Use mfc0 read contents of these registers:
 - » **BadVAddr (register 8)**: contains memory address at which memory reference error occurred
 - » **Status (register 12)**: interrupt mask and enable bits
 - » **Cause (register 13)**: the cause of the exception
 - » **EPC (register 14)**: address of the affected instruction



- Status Register fields:
 - Mask: Interrupt enable
 - » 1 bit for each of 5 hardware and 3 software interrupts
 - k = kernel/user: 0⇒kernel mode
 - e = interrupt enable: 0⇒interrupts disabled
 - **Exception⇒6 LSB shifted left 2 bits, setting 2 LSB to 0:**
 - » run in kernel mode with interrupts disabled

Intel x86 Special Registers



RPL = Requestor Privilege Level

TL = Table Indicator:

(0 = GDT, 1 = LDT)



Index = Index into table

Protected Mode segment selector:


Typical Segment Register
Current Priority is RPL
Of Code Segment (CS)

80386 Special Registers

Segment registers



PG=Paging Enable
 ET=Emulation Type
 TS=Task Switched
 EM=Emulate Coprocessor
 MP=Math coprocessor present
 PE=Protected Mode enable



X=Reserved
 NT=Nested Task
 IOPL=I/O Privilege Level
 OF=Overflow Flag
 DF=Direction Flag
 IF=Interrupt Flag
 TF=Trap Flag
 SF=Sign Flag
 ZF=Zero Flag
 AF=Auxiliary Flag
 PF=Parity Flag
 CF=Carry Flag

Communication



- Now that we have isolated processes, how can they communicate?
 - Shared memory: common mapping to physical page
 - » As long as place objects in shared memory address range, threads from each process can communicate
 - » Note that processes A and B can talk to shared memory through different addresses
 - » In some sense, this violates the whole notion of protection that we have been developing
 - If address spaces don't share memory, all inter-address space communication must go through kernel (via system calls)
 - » Byte stream producer/consumer (put/get): Example, communicate through pipes connecting `stdin/stdout`
 - » Message passing (send/receive): Will explain later how you can use this to build remote procedure call (RPC) abstraction so that you can have one program make procedure calls to another
 - » File System (read/write): File system is shared state!

Closing thought: Protection without Hardware

- Does protection require hardware support for translation and dual-mode behavior?
 - No: Normally use hardware, but anything you can do in hardware can also do in software (possibly expensive)
- Protection via Strong Typing
 - Restrict programming language so that you can't express program that would trash another program
 - Loader needs to make sure that program produced by valid compiler or all bets are off
 - Example languages: LISP, Ada, Modula-3 and Java
- Protection via software fault isolation:
 - Language independent approach: have compiler generate object code that provably can't step out of bounds
 - » Compiler puts in checks for every "dangerous" operation (loads, stores, etc). Again, need special loader.
 - » Alternative, compiler generates "proof" that code cannot do certain things (Proof Carrying Code)
 - Or: use virtual machine to guarantee safe behavior (loads and stores recompiled on fly to check bounds)

Summary

- **Memory is a resource that must be shared**
 - **Controlled Overlap:** only shared when appropriate
 - **Translation:** Change Virtual Addresses into Physical Addresses
 - **Protection:** Prevent unauthorized Sharing of resources
- **Simple Protection through Segmentation**
 - **Base+limit registers** restrict memory accessible to user
 - **Can be used to translate as well**
- **Full translation of addresses through Memory Management Unit (MMU)**
 - **Every Access** translated through page table
 - **Changing of page tables** only available to user

Summary (2/2)

- **Dual-Mode**
 - Kernel/User distinction: User restricted
 - User→Kernel: System calls, Traps, or Interrupts
 - Inter-process communication: shared memory, or through kernel (system calls)
- **Exceptions**
 - Synchronous Exceptions: Traps (including system calls)
 - Asynchronous Exceptions: Interrupts